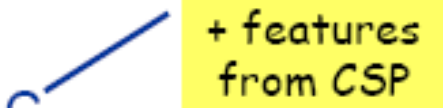


SPIN - Introduction (1)

- **SPIN** (= Simple Promela Interpreter)
 - = is a tool for analysing the logical consistency of **concurrent systems**, specifically of **data communication protocols**.
 - = **state-of-the-art** model checker, used by **>2000 users**
 - Concurrent systems are described in the **modelling language** called **Promela**.
- **Promela** (= Protocol/Process Meta Language)
 - allows for the **dynamic creation** of **concurrent processes**.
 - communication via **message channels** can be defined to be
 - **synchronous** (i.e. rendezvous), or
 - **asynchronous** (i.e. buffered).
 - resembles the programming language **C** 
 - specification language to model **finite-state systems**

SPIN - Introduction (2)

- Major versions:

| | | |
|-----|-----------|-------------------------------------|
| 1.0 | Jan 1991 | initial version [Holzmann 1991] |
| 2.0 | Jan 1995 | partial order reduction |
| 3.0 | Apr 1997 | minimised automaton representation |
| 4.0 | late 2002 | Ax: automata extraction from C code |

- Some success factors of SPIN (subjective!):
 - “press on the button” verification (model checker)
 - very efficient implementation (using C)
 - nice graphical user interface (Xspin)
 - not just a research tool, but well supported
 - contains more than two decades research on advanced computer aided verification (many optimization algorithms)

Documentation on SPIN

- SPIN's starting page:

<http://netlib.bell-labs.com/netlib/spin/whatispin.html>

- Basic SPIN manual
- Getting started with Xspin
- Getting started with SPIN
- Examples and Exercises
- Concise Promela Reference (by Rob Gerth)
- Proceedings of all SPIN Workshops

Also part of SPIN's
documentation distribution
(file: `html.tar.gz`)

- Gerard Holzmann's website for papers on SPIN:

<http://cm.bell-labs.com/cm/cs/who/gerard/>

- SPIN version 1.0 is described in [Holzmann 1991].

Promela Model

- **Promela model** consist of:
 - **type** declarations
 - **channel** declarations
 - **variable** declarations
 - **process** declarations
 - [**init** process]
- A Promela model corresponds with a (usually **very large**, but) **finite transition system**, so
 - no unbounded **data**
 - no unbounded **channels**
 - no unbounded **processes**
 - no unbounded **process creation**

```
mtype = {MSG, ACK};  
chan toS = ...  
chan toR = ...  
bool flag;  
  
proctype Sender() {  
    ...  
} process body  
  
proctype Receiver() {  
    ...  
}  
  
init {  
    ...  
} creates processes
```

Processes (1)

- A **process type** (**proctype**) consist of
 - a **name**
 - a list of **formal parameters**
 - **local variable** declarations
 - **body**

```
proctype Sender(chan in; chan out) {  
    bit sndB, rcvB; local variables  
    do  
        :: out ! MSG, sndB ->  
           in ? ACK, rcvB;  
        if  
            :: sndB == rcvB -> sndB = 1-sndB  
            :: else -> skip  
        fi  
    od  
}
```

body

The body consist of a sequence of **statements**.

Processes (2)

- A **process**
 - is defined by a **proctype** definition
 - executes **concurrently** with all other processes, independent of speed of behaviour
 - **communicate** with other processes
 - using **global** (shared) **variables**
 - using **channels**
- There may be **several processes** of the **same type**.
- Each process has its own **local state**:
 - **process counter** (location within the **proctype**)
 - contents of the **local variables**

Processes (3)

- Process are **created** using the **run** statement (which returns the **process id**).
- Processes can be created at **any point** in the execution (within any process).
- Processes start executing **after** the **run** statement.
- Processes can **also** be created by adding **active** in front of the **proctype** declaration.

```
proctype Foo(byte x) {  
    ...  
}
```

```
init {  
    int pid2 = run Foo(2);  
    run Foo(27);  
}
```

number of procs. (opt.)

```
active[3] proctype Bar() {  
    ...  
}
```

parameters will be
initialised to 0

DEMO

Hello World!

```
/* A "Hello World" Promela model for SPIN. */
active proctype Hello() {
    printf("Hello process, my pid is: %d\n", _pid);
}
init {
    int lastpid;
    printf("init process, my pid is: %d\n", _pid);
    lastpid = run Hello();
    printf("last pid was: %d\n", lastpid);
}
```

random seed

```
$ spin -n2 hello.pr
```

```
init process, my pid is: 1
```

```
    last pid was: 2
```

```
Hello process, my pid is: 0
```

```
    Hello process, my pid is: 2
```

```
3 processes created
```

running SPIN in
random simulation mode

Variables and Types (1)

- Five different (integer) **basic types**.
- **Arrays**
- **Records** (structs)
- **Type conflicts** are detected at runtime.
- **Default initial value** of basic variables (local and global) is **0**.

Basic types

| | | |
|--------------------|-----------------------|---|
| <code>bit</code> | <code>turn=1;</code> | <code>[0..1]</code> |
| <code>bool</code> | <code>flag;</code> | <code>[0..1]</code> |
| <code>byte</code> | <code>counter;</code> | <code>[0..255]</code> |
| <code>short</code> | <code>s;</code> | <code>[-2¹⁶-1.. 2¹⁶-1]</code> |
| <code>int</code> | <code>msg;</code> | <code>[-2³²-1.. 2³²-1]</code> |

Arrays

```
byte a[27];  
bit  flags[4];
```

array
indicing
start at 0

Typedef (records)

```
typedef Record {  
    short f1;  
    byte  f2;  
}  
Record rr;  
rr.f1 = ..
```

variable
declaration

Variables and Types (2)

- Variables should be **declared**.
- Variables can be **given a value** by:
 - **assignment**
 - **argument passing**
 - **message passing**
(see **communication**)
- Variables can be used in **expressions**.

Most **arithmetic**, **relational**, and **logical** operators of C/Java are supported, including **bitshift** operators.

```
int ii;  
bit bb;
```

```
bb=1;  
ii=2;
```

assignment =

```
short s=-1;
```

declaration +
initialisation

```
typedef Foo {  
    bit bb;  
    int ii;  
};
```

```
Foo f;  
f.bb = 0;  
f.ii = -2;
```

equal test ==

```
ii*s+27 == 23;  
printf("value: %d", s*s);
```

Statements (1)

- The body of a process consists of a **sequence of statements**. A statement is either
 - executable**: the statement can be executed **immediately**.
 - blocked**: the statement **cannot** be executed.

executable/blocked
depends on the **global state** of the system.

- An **assignment** is **always executable**.
- An **expression** is also a statement; it is **executable** if it evaluates to **non-zero**.

$2 < 3$

always executable

$x < 27$

only executable if value of x is smaller **27**

$3 + x$

executable if x is not equal to **-3**

Statements (2)

Statements are separated by a semi-colon: ";".

- The **skip** statement is **always executable**.
 - “does nothing”, only changes process' process counter
- A **run** statement is **only executable** if a new process can be created (remember: the number of processes is bounded).
- A **printf** statement is **always executable** (but is not evaluated during verification, of course).

```
int x;  
proctype Aap()  
{  
    int y=1;  
    skip;  
    run Noot();  
    x=2;  
    x>2 && y==1;  
    skip;  
}
```

Executable if **Noot** can be created...

Can only become executable if a **some other process** makes x greater than **2**.



Statements (3)

- **assert**(**<expr>**) ;
 - The **assert**-statement is **always executable**.
 - If **<expr>** evaluates to zero, SPIN will exit with an **error**, as the **<expr>** “**has been violated**”.
 - The **assert**-statement is often used within Promela models, to check whether certain **properties are valid** in a state.

```
proctype monitor() {  
    assert(n <= 3);  
}  
  
proctype receiver() {  
    ...  
    toReceiver ? msg;  
    assert(msg != ERROR);  
    ...  
}
```

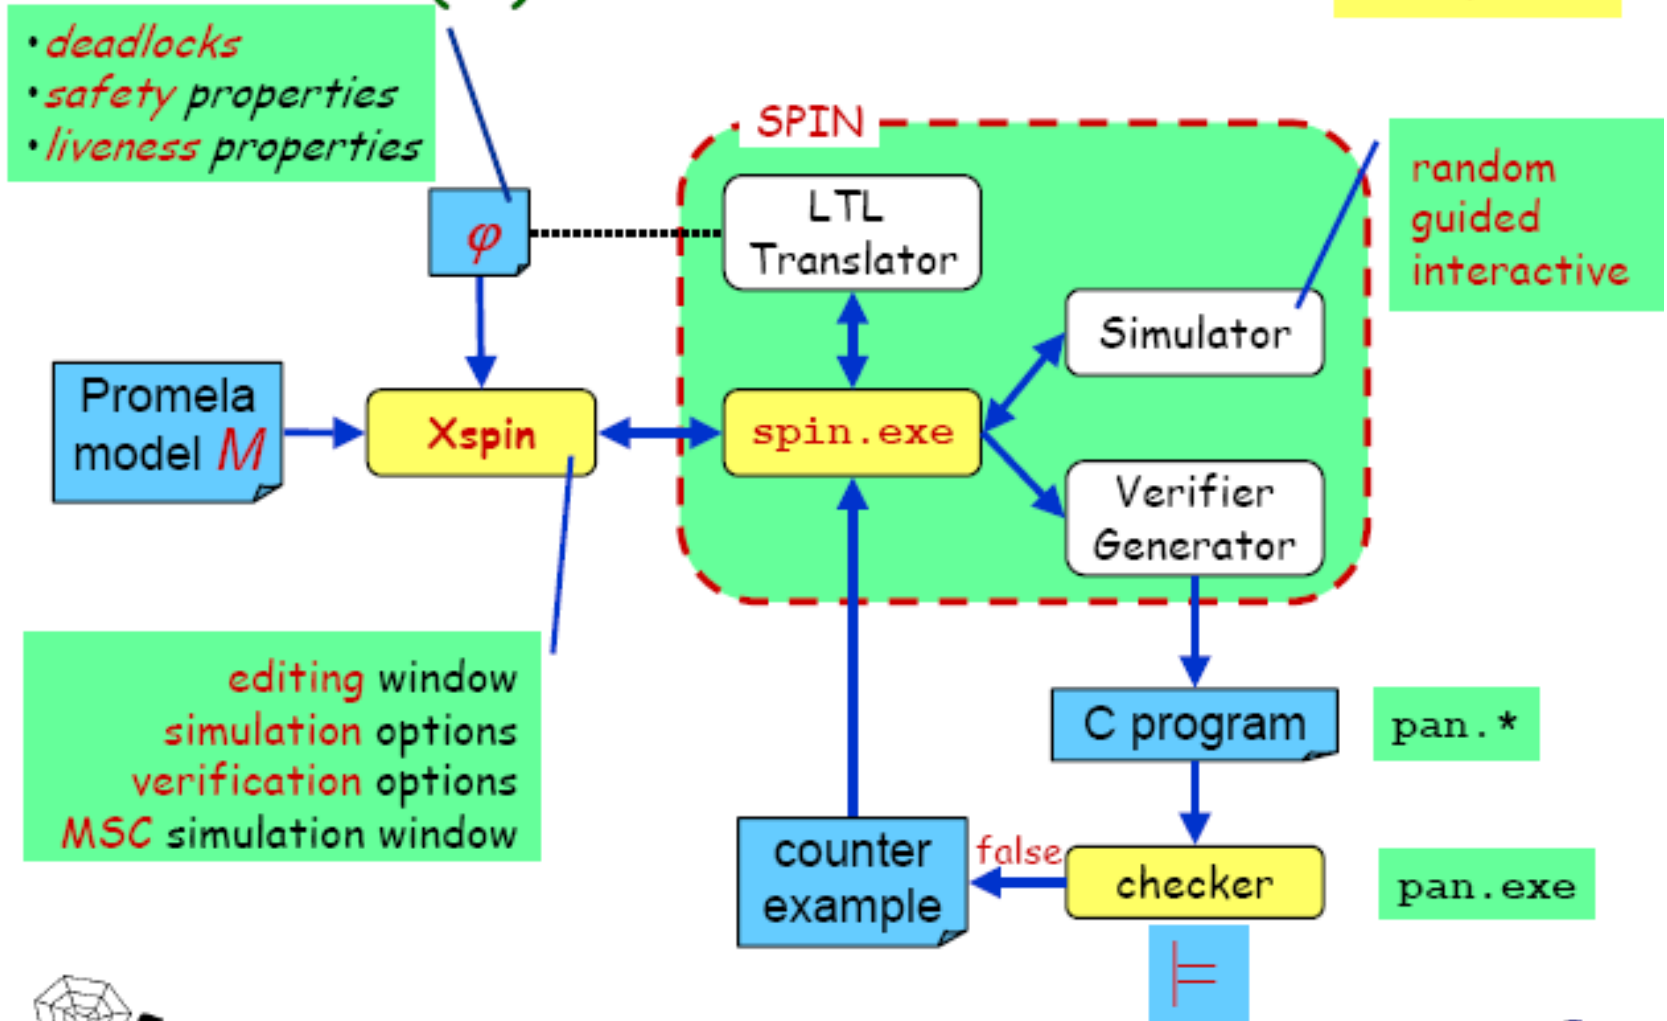
Interleaving Semantics

- Promela **processes** execute **concurrently**.
- **Non-deterministic scheduling** of the processes.
- Processes are **interleaved** (statements of different processes do not occur at the same time).
 - exception: **rendez-vous communication**.
- All statements are **atomic**; each statement is executed without interleaving with other processes.
- Each process may have several **different possible actions** enabled at each point of execution.
 - only one choice is made, **non-deterministically**.

= randomly

(X)SPIN Architecture

$$M \models \varphi$$



Xspin in a nutshell

- **Xspin** allows the user to
 - **edit** Promela models (+ syntax check)
 - **simulate** Promela models
 - random
 - interactive
 - guided
 - **verify** Promela models
 - exhaustive
 - bitstate hashing mode
 - additional **features**
 - Xspin suggest **abstractions** to a Promela model (**slicing**)
 - Xspin can **draw automata** for each process
 - **LTL property manager**
 - **Help** system (with verification/simulation **guidelines**)
- with dialog boxes to set various **options** and **directives** to **tune** the verification process

Mutual Exclusion (1)

```
bit flag;      /* signal entering/leaving the section */
byte mutex;    /* # procs in the critical section.    */

proctype P(bit i) {
    flag != 1;
    flag = 1;
    mutex++;
    printf("MSC: P(%d) has entered section.\n", i);
    mutex--;
    flag = 0;
}

proctype monitor() {
    assert(mutex != 2);
}

init {
    atomic { run P(0); run P(1); run monitor(); }
}
```

models:

```
while (flag == 1) /* wait */;
```

Problem: **assertion violation!**

Both processes can pass the
`flag != 1` "at the same time",
i.e. before `flag` is set to 1.

starts **two** instances of process P

Mutual Exclusion (2)

```
bit  x, y;    /* signal entering/leaving the section */
byte mutex;   /* # of procs in the critical section. */
```

```
active proctype A() {
```

```
  x = 1;
```

```
  y == 0;
```

```
  mutex++;
```

```
  mutex--;
```

```
  x = 0;
```

```
}
```

```
active proctype monitor() {
```

```
  assert(mutex != 2);
```

```
}
```

Process A waits for
process B to end.

```
active proctype B() {
```

```
  y = 1;
```

```
  x == 0;
```

```
  mutex++;
```

```
  mutex--;
```

```
  y = 0;
```

```
}
```

Problem: **invalid-end-state!**

Both processes can pass execute
 $x = 1$ and $y = 1$ "at the same time",
and will then be waiting for each other.



Mutual Exclusion (3)

```
bit  x, y;      /* signal entering/leaving the section */
byte mutex;     /* # of procs in the critical section. */
byte turn;      /* who's turn is it? */

active proctype A() {
    x = 1;
    turn = B_TURN;
    y == 0 ||
        (turn == A_TURN);
    mutex++;
    mutex--;
    x = 0;
}

active proctype B() {
    y = 1;
    turn = A_TURN;
    x == 0 ||
        (turn == B_TURN);
    mutex++;
    mutex--;
    y = 0;
}

active proctype monitor() {
    assert(mutex != 2);
}
```

Can be generalised
to a single process.

First "software-only" solution to the
mutex problem (for two processes).

if-statement (1)

inspired by:
Dijkstra's guarded
command language

```
if
:: choice1 -> stat1.1; stat1.2; stat1.3; ...
:: choice2 -> stat2.1; stat2.2; stat2.3; ...
:: ...
:: choicen -> statn.1; statn.2; statn.3; ...
fi;
```

- If there is at least one **choice_i** (guard) executable, the **if**-statement is executable and SPIN **non-deterministically chooses** one of the executable choices.
- If **no choice_i** is executable, the **if**-statement is **blocked**.
- The operator “**->**” is equivalent to “**;**”. By **convention**, it is used within **if**-statements to **separate** the guards from the statements that follow the guards.

if-statement (2)


```
if
:: (n % 2 != 0) -> n=1
:: (n >= 0)      -> n=n-2
:: (n % 3 == 0) -> n=3
:: else         -> skip
fi
```

- The **else** guard becomes **executable** if **none** of the other guards is executable.

give n a random value

```
if
:: skip -> n=0
:: skip -> n=1
:: skip -> n=2
:: skip -> n=3
fi
```

non-deterministic branching



skips are **redundant**, because assignments are themselves **always executable**...



do-statement (1)

```
do
:: choice1 -> stat1.1; stat1.2; stat1.3; ...
:: choice2 -> stat2.1; stat2.2; stat2.3; ...
:: ...
:: choicen -> statn.1; statn.2; statn.3; ...
od;
```

- With respect to the choices, a **do**-statement behaves in the same way as an **if**-statement.
- However, instead of ending the statement at the end of the chosen list of statements, a **do**-statement **repeats the choice selection**.
- The (**always executable**) **break** statement exits a **do**-loop statement and transfers control to the end of the loop.

do-statement (2)

- Example – modelling a traffic light

if- and **do**-statements are ordinary Promela statements; so they can be nested.


```
mtype = { RED, YELLOW, GREEN } ;
```

mtype (message type) models enumerations in Promela

```
active proctype TrafficLight() {  
    byte state = GREEN;  
    do  
        :: (state == GREEN)    -> state = YELLOW;  
        :: (state == YELLOW)  -> state = RED;  
        :: (state == RED)     -> state = GREEN;  
    od;  
}
```

Note: this **do**-loop does not contain any non-deterministic choice.

Communication (1)



Communication (2)

- Communication between processes is via **channels**:
 - **message passing**
 - **rendez-vous** synchronisation (**handshake**)

- Both are defined as **channels**:

also called:
queue or **buffer**

```
chan <name> = [<dim>] of {<t1>, <t2>, ... <tn>} ;
```


name of
the channel

type of the elements that will be
transmitted over the channel

number of elements in the channel
dim==0 is special case: **rendez-vous**

```
chan c      = [1] of {bit};  
chan toR    = [2] of {mtype, bit};  
chan line[2] = [1] of {mtype, Record};
```

array of
channels



Communication (3)

- channel = **FIFO**-buffer (for **dim**>0)

! **Sending** - *putting a message into a channel*

ch ! <expr₁>, <expr₂>, ... <expr_n>;

- The values of **<expr_i>** should correspond with the types of the channel declaration.
- A **send**-statement is **executable** if the channel is **not full**.

? **Receiving** - *getting a message out of a channel*

**<var> +
<const>
can be
mixed**

ch ? <var₁>, <var₂>, ... <var_n>;

message passing

- If the channel is **not empty**, the message is fetched from the channel and the individual parts of the message are stored into the **<var_i>**s.

ch ? <const₁>, <const₂>, ... <const_n>;

message testing

- If the channel is **not empty** and the message at the front of the channel evaluates to the individual **<const_i>**, the statement is executable and the message is removed from the channel.

Communication (4)

- Rendez-vous communication

`<dim> == 0`

The number of elements in the channel is now zero.

- If `send ch!` is enabled and if there is a corresponding `receive ch?` that can be executed simultaneously and the constants match, then both statements are enabled.
- Both statements will “handshake” and together take the transition.

- *Example:*

`chan ch = [0] of {bit, byte};`

- P wants to do `ch ! 1, 3+7`
- Q wants to do `ch ? 1, x`
- Then after the communication, `x` will have the value 10.

Promela Model

- A Promela model consist of:

- **type** declarations

mtype, typedefs,
constants

- **channel** declarations

chan ch = [dim] of {type, ...}
asynchronous: dim > 0
rendez-vous: dim == 0

- **global variable** declarations

can be accessed
by **all** processes

- **process** declarations

behaviour of the processes:
local variables + statements

- [**init** process]

initialises variables and
starts processes



Promela statements

are either **executable**
or **blocked**

| | |
|-------------------------------|--|
| skip | always executable |
| assert (<expr>) | always executable |
| <i>expression</i> | executable if not zero |
| <i>assignment</i> | always executable |
| if | executable if at least one guard is executable |
| do | executable if at least one guard is executable |
| break | always executable (exits do -statement) |
| <i>send</i> (ch !) | executable if channel ch is not full |
| <i>receive</i> (ch ?) | executable if channel ch is not empty |

atomic

```
atomic { stat1; stat2; ... statn }
```

- can be used to group statements into an atomic sequence;
all statements are executed in a single step
(no interleaving with statements of other processes)
 - is executable if `stat1` is executable / no pure atomicity
 - if a `stati` (with `i > 1`) is blocked, the “atomicity token” is
(temporarily) lost and other processes may do a step
- (Hardware) solution to the mutual exclusion problem:

```
proctype P(bit i) {  
    atomic {flag != 1; flag = 1; }  
    mutex++;  
    mutex--;  
    flag = 0;  
}
```



d_step

```
d_step { stat1; stat2; ... statn }
```

- more **efficient** version of **atomic**: **no intermediate states** are generated and stored
- may only contain **deterministic** steps
- it is a **run-time error** if **stat_i** ($i > 1$) blocks.
- **d_step** is especially useful to perform intermediate computations in a **single transition**

```
:: Rout?i(v) -> d_step {  
    k++;  
    e[k].ind = i;  
    e[k].val = v;  
    i=0; v=0 ;  
}
```

- **atomic** and **d_step** can be used to **lower** the number of **states** of the model

timeout (1)

- Promela does **not** have **real-time** features.
 - In Promela we can only specify **functional behaviour**.
 - Most protocols, however, use **timers** or a **timeout** mechanism to **resend** messages or acknowledgements.
- **timeout**
 - SPIN's **timeout** becomes **executable** if there is **no other process** in the system which is executable
 - so, **timeout** models a **global timeout**
 - **timeout** provides an **escape** from **deadlock states**
 - **beware of statements** that are always executable...

timeout (2)

- Example to recover from message loss:

```
active proctype Receiver()
{
    bit rcvbit;
    do
        :: toR ? MSG, rcvbit -> toS ! ACK, rcvbit;
        :: timeout          -> toS ! ACK, rcvbit;
    od
}
```

- Premature timeouts can be modelled by replacing the timeout by skip (which is always executable).

One might want to limit the number of premature timeouts (see [Ruys & Langerak 1997]).



goto

goto label

- transfers execution to label
- each Promela statement might be labelled
- quite useful in modelling communication protocols

```
wait_ack:
  if
  :: B?ACK -> ab=1-ab ; goto success
  :: ChunkTimeout?SHAKE ->
    if
    :: (rc < MAX) -> rc++; F!(i==1), (i==n), ab, d[i];
                      goto wait_ack
    :: (rc >= MAX) -> goto error
    fi
  fi ;
```

Timeout modelled by a channel.

Part of model of BRP

unless

```
{ <stats> } unless { guard; <stats> }
```

- Statements in *<stats>* are executed **until** the first statement (*guard*) in the escape sequence becomes executable.
- resembles **exception handling** in languages like Java
- *Example:*

```
proctype MicroProcessor() {  
    {  
        ...  
        /* execute normal instructions */  
    }  
    unless { port ? INTERRUPT; ... }  
}
```

macros - **cpp** preprocessor

- Promela uses **cpp**, the **C preprocessor** to preprocess Promela models. This is useful to define:

- **constants**

```
#define MAX 4
```

All **cpp** commands start with a **hash**:
#define, **#ifdef**, **#include**, etc.

- **macros**

```
#define RESET_ARRAY(a) \  
    d_step { a[0]=0; a[1]=0; a[2]=0; a[3]=0; }
```

- **conditional** Promela model fragments

```
#define LOSSY 1  
...  
#ifdef LOSSY  
active proctype Daemon() { /* steal messages */ }  
#endif
```



`inline` - poor man's procedures

- Promela also has its own `macro-expansion` feature using the `inline`-construct.

```
inline init_array(a) {  
  d_step {  
    i=0;  
    do  
      :: i<N -> a[i] = 0; i++  
      :: else -> break  
    od;  
    i=0;  
  }  
}
```

Should be `declared somewhere else` (probably as a local variable).

Be sure to `reset` temporary variables.

- error messages are more `useful` than when using `#define`
- `cannot` be used as `expression`
- all `variables` should be `declared somewhere else`

State vector

- A **state vector** is the information to uniquely identify a **system state**; it contains:
 - **global variables**
 - contents of the **channels**
 - for each **process** in the system:
 - **local variables**
 - **process counter** of the process
- It is important to **minimise** the **size** of the **state vector**.

state vector = m bytes
state space = n states



storing the state space
may require $n*m$ bytes

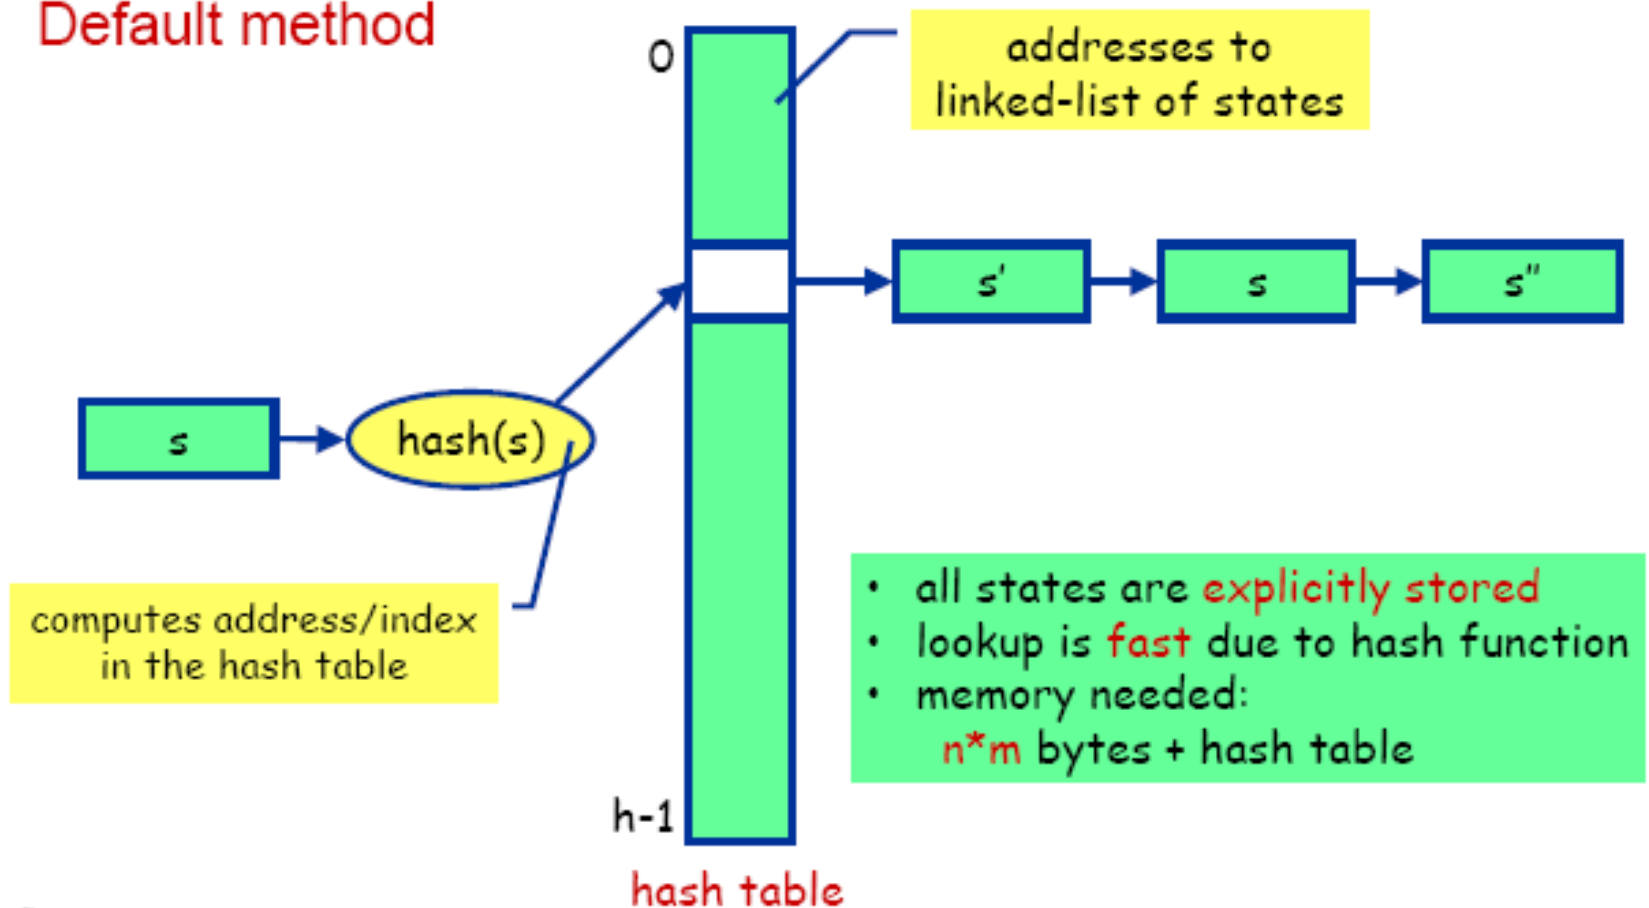
SPIN provides several algorithms to
compress the state vector.

[Holzmann 1997 - State Compression]



Storing States in SPIN

Default method



SPIN Verification Report

(Spin Version 3.4.12 -- 18 December 2001)

the size of a single state on

Full statespace search for:

never-claim

assertion violations

cycle checks

invalid endstates

- (not selected)

+

- (disabled by -DSAFETY)

+

longest execution path

State-vector 96 byte, depth reached 18637, errors: 0

169208 states, stored

71378 states, matched

240586 transitions (= stored+matched)

31120 atomic steps

hash conflicts: 150999 (resolved)

(max size 2¹⁹ states)

property was
satisfied

total number of states
(i.e. the state space)

Stats on memory usage (in Megabytes):

17.598 equivalent memory usage for states

(stored*(State-vector + overhead))

11.634 actual memory usage for states (compression: 66.11%)

State-vector as stored = 61 byte + 8 byte overhead

2.097 memory used for hash-table (-w19)

0.480 memory used for DFS stack (-m20000)

14.354 total actual memory usage

total amount of memory used for this verification



A tutorial by Theo Ruys